

Search Engines

Information Retrieval in Practice


Indexes

- *Indexes* are data structures designed to make search faster
- Text search has unique requirements, which leads to unique data structures
- Most common data structure is *inverted index*
 - general name for a class of structures
 - “inverted” because documents are associated with words, rather than words with documents
 - similar to a *concordance*

Indexes and Ranking

- Indexes are designed to support *search*
 - faster response time, supports updates
- Text search engines use a particular form of search: *ranking*
 - documents are retrieved in sorted order according to a score computing using the document representation, the query, and a *ranking algorithm*
- What is a reasonable abstract model for ranking?
 - enables discussion of indexes without details of retrieval model

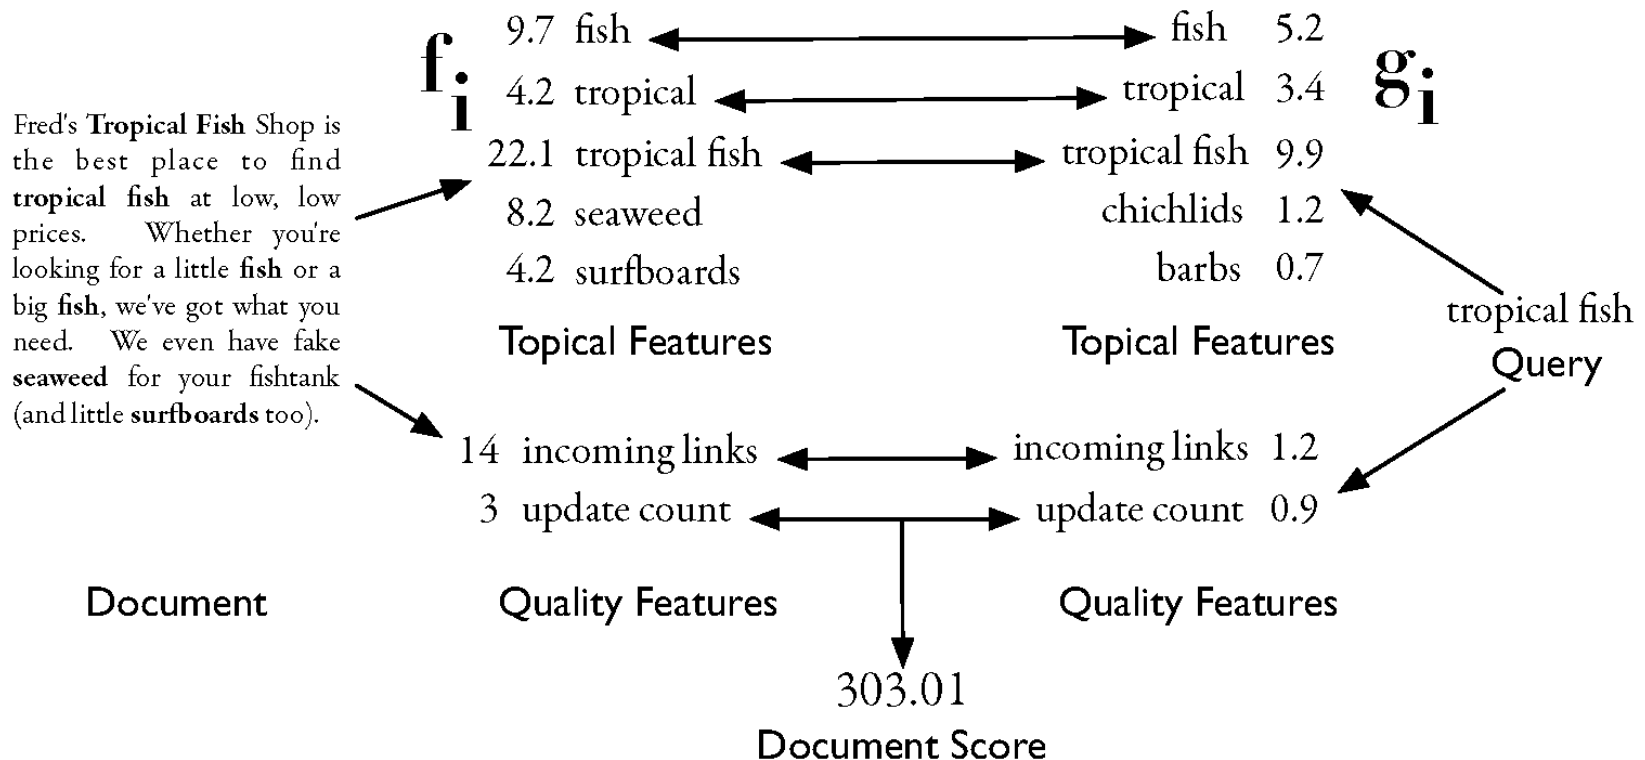
Abstract Model of Ranking



More Concrete Model

$$R(Q, D) = \sum_i g_i(Q) f_i(D)$$

f_i is a document feature function
 g_i is a query feature function



Inverted Index

- Each index term is associated with an *inverted list*
 - Contains lists of documents, or lists of word occurrences in documents, and other information
 - Each entry is called a *posting*
 - The part of the posting that refers to a specific document or location is called a *pointer*
 - Each document in the collection is given a unique number
 - Lists are usually *document-ordered* (sorted by document number)

Example “Collection”

- S_1 Tropical fish include fish found in tropical environments around the world, including both freshwater and salt water species.
- S_2 Fishkeepers often use the term tropical fish to refer only those requiring fresh water, with saltwater tropical fish referred to as marine fish.
- S_3 Tropical fish are popular aquarium fish, due to their often bright coloration.
- S_4 In freshwater fish, this coloration typically derives from iridescence, while salt water fish are generally pigmented.

Four sentences from the Wikipedia entry for *tropical fish*

Simple Inverted Index

and	1				only	2			
aquarium	3				pigmented	4			
are	3	4			popular	3			
around	1				refer	2			
as	2				referred	2			
both	1				requiring	2			
bright	3				salt	1	4		
coloration	3	4			saltwater	2			
derives	4				species	1			
due	3				term	2			
environments	1				the	1	2		
fish	1	2	3	4	their	3			
fishkeepers	2				this	4			
found	1				those	2			
fresh	2				to	2	3		
freshwater	1	4			tropical	1	2	3	
from	4				typically	4			
generally	4				use	2			
in	1	4			water	1	2	4	
include	1				while	4			
including	1				with	2			
iridescence	4				world	1			
marine	2								
often	2	3							

Inverted Index with counts

- supports better ranking algorithms

and	1:1				only	2:1			
aquarium	3:1				pigmented	4:1			
are	3:1	4:1			popular	3:1			
around	1:1				refer	2:1			
as	2:1				referred	2:1			
both	1:1				requiring	2:1			
bright	3:1				salt	1:1	4:1		
coloration	3:1	4:1			saltwater	2:1			
derives	4:1				species	1:1			
due	3:1				term	2:1			
environments	1:1				the	1:1	2:1		
fish	1:2	2:3	3:2	4:2	their	3:1			
fishkeepers	2:1				this	4:1			
found	1:1				those	2:1			
fresh	2:1				to	2:2	3:1		
freshwater	1:1	4:1			tropical	1:2	2:2	3:1	
from	4:1				typically	4:1			
generally	4:1				use	2:1			
in	1:1	4:1			water	1:1	2:1	4:1	
include	1:1				while	4:1			
including	1:1				with	2:1			
iridescence	4:1				world	1:1			
marine	2:1								
often	2:1	3:1							

- supports proximity matches

[illegible]

Proximity Matches

- Matching phrases or words within a window
 - e.g., "tropical fish", or "find tropical within 5 words of fish"
- Word positions in inverted lists make these types of query features efficient
 - e.g.,


tropical	1,1		1,7	2,6	2,17		3,1			
fish	1,2	1,4		2,7	2,18	2,23	3,2	3,6	4,3	4,13

Fields and Extents

- Document structure is useful in search
 - *field* restrictions
 - e.g., date, from:, etc.
 - some fields more important
 - e.g., title
- Options:
 - separate inverted lists for each field type
 - add information about fields to postings
 - use *extent lists*

Extent Lists

- An *extent* is a contiguous region of a document
 - represent extents using word positions
 - inverted list records all extents for a given field type
 - e.g.,



Other Issues

- Precomputed scores in inverted list
 - e.g., list for “fish” [(1:3.6), (3:2.2)], where 3.6 is total feature value for document 1
 - improves speed but reduces flexibility
- Score-ordered lists
 - query processing engine can focus only on the top part of each inverted list, where the highest-scoring documents are recorded
 - very efficient for single-word queries

Compression

- Inverted lists are very large
 - e.g., 25-50% of collection for TREC collections using Indri search engine
 - Much higher if n-grams are indexed
- Compression of indexes saves disk and/or memory space
 - Typically have to decompress lists to use them
 - Best compression techniques have good *compression ratios* and are easy to decompress
- *Lossless* compression – no information lost

Compression

- *Basic idea*: Common data elements use short codes while uncommon data elements use longer codes
 - Example: coding numbers
 - number sequence:
0, 1, 0, 3, 0, 2, 0
 - possible encoding:
00 01 00 10 00 11 00
 - encode 0 using a single 0:
0 01 0 10 0 11 0
 - only 10 bits, but...

Compression Example

- *Ambiguous* encoding – not clear how to decode

- another decoding:

0 01 01 0 0 11 0

- which represents:

0, 1, 1, 0, 0, 3, 0

- use unambiguous code:

Number	Code
0	0
1	101
2	110
3	111

- which gives:

0 101 0 111 0 110 0

Delta Encoding

- Word count data is good candidate for compression
 - many small numbers and few larger numbers
 - encode small numbers with small codes
- Document numbers are less predictable
 - but differences between numbers in an ordered list are smaller and more predictable
- *Delta encoding*:
 - encoding differences between document numbers (*d-gaps*)

Delta Encoding

- Inverted list (without counts)

1, 5, 9, 18, 23, 24, 30, 44, 45, 48

- Differences between adjacent numbers

1, 4, 4, 9, 5, 1, 6, 14, 1, 3

- Differences for a high-frequency word are easier to compress, e.g.,

1, 1, 2, 1, 5, 1, 4, 1, 1, 3, ...

- Differences for a low-frequency word are large, e.g.,

109, 3766, 453, 1867, 992, ...

Bit-Aligned Codes

- Breaks between encoded numbers can occur after any bit position
- *Unary* code
 - Encode k by k 1s followed by 0
 - 0 at end makes code unambiguous

Number	Code
0	0
1	10
2	110
3	1110
4	11110
5	111110

Unary and Binary Codes

- Unary is very efficient for small numbers such as 0 and 1, but quickly becomes very expensive
 - 1023 can be represented in 10 binary bits, but requires 1024 bits in unary
- Binary is more efficient for large numbers, but it may be ambiguous

Elias- γ Code

- To encode a number k , compute
 - $k_d = \lfloor \log_2 k \rfloor$
 - $k_r = k - 2^{\lfloor \log_2 k \rfloor}$
- k_d is number of binary digits, encoded in unary

Number (k)	k_d	k_r	Code
1	0	0	0
2	1	0	10 0
3	1	1	10 1
6	2	2	110 10
15	3	7	1110 111
16	4	0	11110 0000
255	7	127	11111110 1111111
1023	9	511	1111111110 111111111

Elias- δ Code

- Elias- γ code uses no more bits than unary, many fewer for $k > 2$
 - 1023 takes 19 bits instead of 1024 bits using unary
- In general, takes $2\lfloor \log_2 k \rfloor + 1$ bits
- To improve coding of large numbers, use Elias- δ code
 - Instead of encoding k_d in unary, we encode $k_d + 1$ using Elias- γ
 - Takes approximately $2 \log_2 \log_2 k + \log_2 k$ bits

Elias- δ Code

- Split k_d into:
 - $k_{dd} = \lfloor \log_2(k_d + 1) \rfloor$
 - $k_{dr} = k_d - 2^{\lfloor \log_2(k_d + 1) \rfloor}$
- encode k_{dd} in unary, k_{dr} in binary, and k_r in binary

Number (k)	k_d	k_r	k_{dd}	k_{dr}	Code
1	0	0	0	0	0
2	1	0	1	0	10 0 0
3	1	1	1	0	10 0 1
6	2	2	1	1	10 1 10
15	3	7	2	0	110 00 111
16	4	0	2	1	110 01 0000
255	7	127	3	0	1110 000 1111111
1023	9	511	3	2	1110 010 111111111


```

#
# Generating Elias-gamma and Elias-delta codes in Python
#

import math

def unary_encode(n):
    return "1" * n + "0"

def binary_encode(n, width):
    r = ""
    for i in range(0,width):
        if ((1<<i) & n) > 0:
            r = "1" + r
        else:
            r = "0" + r
    return r

def gamma_encode(n):
    logn = int(math.log(n,2))
    return unary_encode( logn ) + " " + binary_encode(n, logn)

def delta_encode(n):
    logn = int(math.log(n,2))
if n == 1:
    return "0"
else:
    loglog = int(math.log(logn+1,2))
    residual = logn+1 - int(math.pow(2, loglog))
    return unary_encode( loglog ) + " " + binary_encode( residual, loglog ) + " " + binary_encode(n, logn)

if __name__ == "__main__":
    for n in [1,2,3, 6, 15,16,255,1023]:
        logn = int(math.log(n,2))
        loglogn = int(math.log(logn+1,2))
        print n, "d_r", logn
        print n, "d_dd", loglogn
        print n, "d_dr", logn + 1 - int(math.pow(2,loglogn))
        print n, "delta", delta_encode(n)
        #print n, "gamma", gamma_encode(n)
        #print n, "binary", binary_encode(n)

```

Byte-Aligned Codes

- Variable-length bit encodings can be a problem on processors that process bytes
- *v-byte* is a popular byte-aligned code
 - Similar to Unicode UTF-8
- Shortest v-byte code is 1 byte
- Numbers are 1 to 4 bytes, with high bit 1 in the last byte, 0 otherwise

V-Byte Encoding

k	Number of bytes
$k < 2^7$	1
$2^7 \leq k < 2^{14}$	2
$2^{14} \leq k < 2^{21}$	3
$2^{21} \leq k < 2^{28}$	4

k	Binary Code	Hexadecimal
1	1 0000001	81
6	1 0000110	86
127	1 1111111	FF
128	0 0000001 1 0000000	01 80
130	0 0000001 1 0000010	01 82
20000	0 0000001 0 0011100 1 0100000	01 1C A0

V-Byte Encoder

```
public void encode( int[] input, ByteBuffer output ) {  
    for( int i : input ) {  
        while( i >= 128 ) {  
            output.put( i & 0x7F );  
            i >>>= 7;  
        }  
        output.put( i | 0x80 );  
    }  
}
```

V-Byte Decoder

```
public void decode( byte[] input, IntBuffer output ) {  
    for( int i=0; i < input.length; i++ ) {  
        int position = 0;  
        int result = ((int)input[i] & 0x7F);  
  
        while( (input[i] & 0x80) == 0 ) {  
            i += 1;  
            position += 1;  
            int unsignedByte = ((int)input[i] & 0x7F);  
            result |= (unsignedByte << (7*position));  
        }  
  
        output.put(result);  
    }  
}
```

Compression Example

- Consider invert list with positions:

$(1, 2, [1, 7])(2, 3, [6, 17, 197])(3, 1, [1])$

- Delta encode document numbers and positions:

$(1, 2, [1, 6])(1, 3, [6, 11, 180])(1, 1, [1])$

- Compress using v-byte:


81 82 81 86 81 82 86 8B 01 B4 81 81 81

Skipping

- Search involves comparison of inverted lists of different lengths
 - Can be very inefficient
 - “Skipping” ahead to check document numbers is much better
 - Compression makes this difficult
 - Variable size, only d-gaps stored
- Skip pointers are additional data structure to support skipping

Skip Pointers

- A skip pointer (d, p) contains a document number d and a byte (or bit) position p
 - Means there is an inverted list posting that starts at position p , and the posting before it was for document d



Skip Pointers

- Example

- Inverted list

5, 11, 17, 21, 26, 34, 36, 37, 45, 48, 51, 52, 57, 80, 89, 91, 94, 101, 104, 119

- D-gaps

5, 6, 6, 4, 5, 9, 2, 1, 8, 3, 3, 1, 5, 23, 9, 2, 3, 7, 3, 15

- Skip pointers

(17, 3), (34, 6), (45, 9), (52, 12), (89, 15), (101, 18)

Auxiliary Structures

- Inverted lists usually stored together in a single file for efficiency
 - *Inverted file*
- *Vocabulary or lexicon*
 - Contains a lookup table from index terms to the byte offset of the inverted list in the inverted file
 - Either hash table in memory or B-tree for larger vocabularies
- Term statistics stored at start of inverted lists
- Collection statistics stored in separate file

Index Construction

- Simple in-memory indexer

procedure BUILDINDEX(D)

$I \leftarrow \text{HashTable}()$

$n \leftarrow 0$

for all documents $d \in D$ **do**

$n \leftarrow n + 1$

$T \leftarrow \text{Parse}(d)$

 Remove duplicates from T

for all tokens $t \in T$ **do**

if $I_t \notin I$ **then**

$I_t \leftarrow \text{Array}()$

end if

$I_t.\text{append}(n)$

end for

end for

return I

end procedure

▷ D is a set of text documents

▷ Inverted list storage

▷ Document numbering

▷ Parse document into tokens

Merging

- Merging addresses limited memory problem
 - Build the inverted list structure until memory runs out
 - Then write the partial index to disk, start making a new one
 - At the end of this process, the disk is filled with many partial indexes, which are merged
- Partial lists must be designed so they can be merged in small pieces
 - e.g., storing in alphabetical order

Merging

Index A

aardvark	2	3	4	5	apple	2	4
----------	---	---	---	---	-------	---	---

Index B

aardvark	6	9	actor	15	42	68
----------	---	---	-------	----	----	----

Index A

aardvark	2	3	4	5								apple	2	4
----------	---	---	---	---	--	--	--	--	--	--	--	-------	---	---

Index B

aardvark						6	9	actor	15	42	68			
----------	--	--	--	--	--	---	---	-------	----	----	----	--	--	--

Combined index

aardvark	2	3	4	5	6	9	actor	15	42	68	apple	2	4
----------	---	---	---	---	---	---	-------	----	----	----	-------	---	---

Distributed Indexing

- Distributed processing driven by need to index and analyze huge amounts of data (i.e., the Web)
- Large numbers of inexpensive servers used rather than larger, more expensive machines
- *MapReduce* is a distributed programming tool designed for indexing and analysis tasks

Example

- Given a large text file that contains data about credit card transactions
 - Each line of the file contains a credit card number and an amount of money
 - Determine the number of unique credit card numbers
- Could use hash table – memory problems
 - counting is simple with sorted file
- Similar with distributed approach
 - sorting and placement are crucial

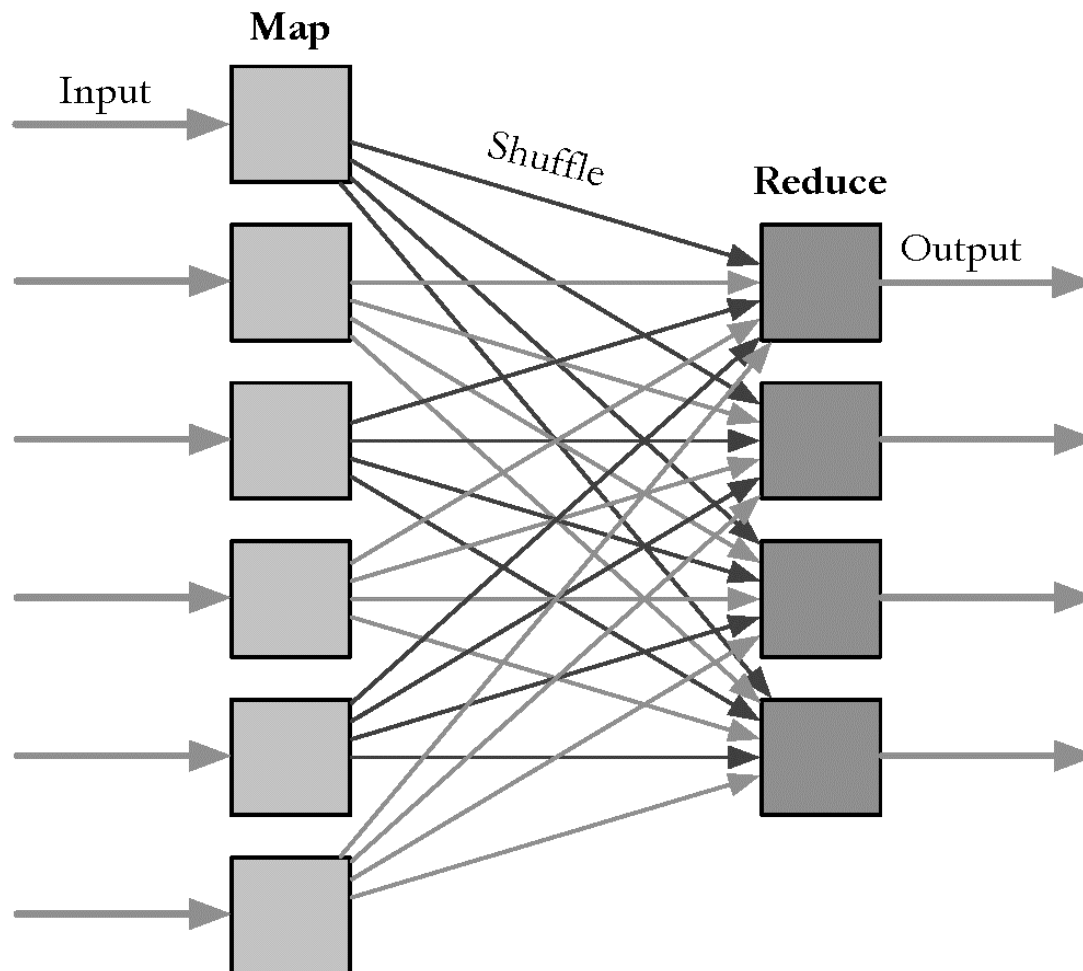
MapReduce

- Distributed programming framework that focuses on data placement and distribution
- *Mapper*
 - Generally, transforms a list of items into another list of items of the same length
- *Reducer*
 - Transforms a list of items into a single item
 - Definitions not so strict in terms of number of outputs
- Many mapper and reducer tasks on a cluster of machines

MapReduce

- Basic process
 - *Map* stage which transforms data records into pairs, each with a key and a value
 - *Shuffle* uses a hash function so that all pairs with the same key end up next to each other and on the same machine
 - *Reduce* stage processes records in batches, where all pairs with the same key are processed at the same time
- *Idempotence* of Mapper and Reducer provides fault tolerance
 - multiple operations on same input gives same output

MapReduce



Example

```
procedure MAPCREDITCARDS(input)
  while not input.done() do
    record  $\leftarrow$  input.next()
    card  $\leftarrow$  record.card
    amount  $\leftarrow$  record.amount
    Emit(card, amount)
  end while
end procedure
```

```
procedure REDUCECREDITCARDS(key, values)
  total  $\leftarrow$  0
  card  $\leftarrow$  key
  while not values.done() do
    amount  $\leftarrow$  values.next()
    total  $\leftarrow$  total + amount
  end while
  Emit(card, total)
end procedure
```

Indexing Example

```
procedure MAPDOCUMENTSTOPOSTINGS(input)
  while not input.done() do
    document  $\leftarrow$  input.next()
    number  $\leftarrow$  document.number
    position  $\leftarrow$  0
    tokens  $\leftarrow$  Parse(document)
    for each word  $w$  in tokens do
      Emit( $w$ , number:position)
      position = position + 1
    end for
  end while
end procedure
```

```
procedure REDUCEPOSTINGSTOLISTS(key, values)
  word  $\leftarrow$  key
  WriteWord(word)
  while not input.done() do
    EncodePosting(values.next())
  end while
end procedure
```

Result Merging

- Index merging is a good strategy for handling updates when they come in large batches
- For small updates this is very inefficient
 - instead, create separate index for new documents, merge *results* from both searches
 - could be in-memory, fast to update and search
- Deletions handled using *delete list*
 - Modifications done by putting old version on delete list, adding new version to new documents index

Query Processing

- Document-at-a-time
 - Calculates complete scores for documents by processing all term lists, one document at a time
- Term-at-a-time
 - Accumulates scores for documents by processing term lists one at a time
- Both approaches have optimization techniques that significantly reduce time required to generate scores

Document-At-A-Time

salt	1:1				4:1
water	1:1	2:1			4:1
tropical	1:2	2:2	3:1		
score	1:4	2:3	3:1		4:2

Pseudocode Function Descriptions

- `getCurrentDocument()`
 - Returns the document number of the current posting of the inverted list.
- `skipForwardToDocument(d)`
 - Moves forward in the inverted list until `getCurrentDocument() <= d`. This function may read to the end of the list.
- `movePastDocument(d)`
 - Moves forward in the inverted list until `getCurrentDocument() < d`.
- `moveToNextDocument()`
 - Moves to the next document in the list. Equivalent to `movePastDocument(getCurrentDocument())`.
- `getNextAccumulator(d)`
 - returns the first document number $d' \geq d$ that has already has an accumulator.
- `removeAccumulatorsBetween(a, b)`
 - Removes all accumulators for documents numbers between a and b . A_d will be removed iff $a < d < b$.

Document-At-A-Time

```
procedure DOCUMENTATATIMEREtrieval( $Q, I, f, g, k$ )  
   $L \leftarrow \text{Array}()$   
   $R \leftarrow \text{PriorityQueue}(k)$   
  for all terms  $w_i$  in  $Q$  do  
     $l_i \leftarrow \text{InvertedList}(w_i, I)$   
     $L.\text{add}( l_i )$   
  end for  
  for all documents  $d \in I$  do  
     $s_d \leftarrow 0$   
    for all inverted lists  $l_i$  in  $L$  do  
      if  $l_i.\text{getCurrentDocument}() = d$  then  
         $s_d \leftarrow s_d + g_i(Q)f_i(l_i)$  ▷ Update the document score  
      end if  
       $l_i.\text{movePastDocument}( d )$   
    end for  
     $R.\text{add}( s_d, d )$   
  end for  
  return the top  $k$  results from  $R$   
end procedure
```

Term-At-A-Time

	salt	1:1	4:1
partial scores		1:1	4:1

old partial scores	1:1		4:1	
	water	1:1	2:1	4:1
new partial scores	1:2	2:1	4:2	

old partial scores	1:2	2:1		4:2
tropical	1:2	2:2	3:1	
final scores	1:4	2:3	2:2	4:2

Term-At-A-Time

```
procedure TERMATATIMEREtrieval( $Q, I, f, g, k$ )  
   $A \leftarrow \text{HashTable}()$   
   $L \leftarrow \text{Array}()$   
   $R \leftarrow \text{PriorityQueue}(k)$   
  for all terms  $w_i$  in  $Q$  do  
     $l_i \leftarrow \text{InvertedList}(w_i, I)$   
     $L.\text{add}(l_i)$   
  end for  
  for all lists  $l_i \in L$  do  
    while  $l_i$  is not finished do  
       $d \leftarrow l_i.\text{getCurrentDocument}()$   
       $A_d \leftarrow A_d + g_i(Q)f(l_i)$   
       $l_i.\text{moveToNextDocument}()$   
    end while  
  end for  
  for all accumulators  $A_d$  in  $A$  do  
     $s_d \leftarrow A_d$   $\triangleright$  Accumulator contains the document score  
     $R.\text{add}(s_d, d)$   
  end for  
  return the top  $k$  results from  $R$   
end procedure
```

Optimization Techniques

- Term-at-a-time uses more memory for accumulators, but accesses disk more efficiently
- Two classes of optimization
 - Read less data from inverted lists
 - e.g., skip lists
 - better for simple feature functions
 - Calculate scores for fewer documents
 - e.g., conjunctive processing
 - better for complex feature functions

```

1: procedure TERMATATIMERETRIEVAL( $Q, I, f, g, k$ )
2:    $A \leftarrow \text{Map}()$ 
3:    $L \leftarrow \text{Array}()$ 
4:    $R \leftarrow \text{PriorityQueue}(k)$ 
5:   for all terms  $w_i$  in  $Q$  do
6:      $l_i \leftarrow \text{InvertedList}(w_i, I)$ 
7:      $L.\text{add}(l_i)$ 
8:   end for
9:   for all lists  $l_i \in L$  do
10:     $d_0 \leftarrow -1$ 
11:    while  $l_i$  is not finished do
12:      if  $i = 0$  then
13:         $d \leftarrow l_i.\text{getCurrentDocument}()$ 
14:         $A_d \leftarrow A_d + g_i(Q)f(l_i)$ 
15:         $l_i.\text{moveToNextDocument}()$ 
16:      else
17:         $d \leftarrow l_i.\text{getCurrentDocument}()$ 
18:         $d' \leftarrow A.\text{getNextAccumulator}(d)$ 
19:         $A.\text{removeAccumulatorsBetween}(d_0, d')$ 
20:        if  $d = d'$  then
21:           $A_d \leftarrow A_d + g_i(Q)f(l_i)$ 
22:           $l_i.\text{moveToNextDocument}()$ 
23:        else
24:           $l_i.\text{skipForwardToDocument}(d')$ 
25:        end if
26:         $d_0 \leftarrow d'$ 
27:      end if
28:    end while
29:  end for
30:  for all accumulators  $A_d$  in  $A$  do
31:     $s_d \leftarrow A_d$   $\triangleright$  Accumulator contains the document score
32:     $R.\text{add}(s_d, d)$ 
33:  end for
34:  return the top  $k$  results from  $R$ 
35: end procedure

```

Conjunctive Term-at-a-Time

```

1: procedure DOCUMENTATATIMERETRIEVAL( $Q, I, f, g, k$ )
2:    $L \leftarrow \text{Array}()$ 
3:    $R \leftarrow \text{PriorityQueue}(k)$ 
4:   for all terms  $w_i$  in  $Q$  do
5:      $l_i \leftarrow \text{InvertedList}(w_i, I)$ 
6:      $L.\text{add}( l_i )$ 
7:   end for
8:    $d \leftarrow -1$ 
9:   while all lists in  $L$  are not finished do
10:     $s_d \leftarrow 0$ 
11:    for all inverted lists  $l_i$  in  $L$  do
12:      if  $l_i.\text{getCurrentDocument}() > d$  then
13:         $d \leftarrow l_i.\text{getCurrentDocument}()$ 
14:      end if
15:    end for
16:    for all inverted lists  $l_i$  in  $L$  do
17:       $l_i.\text{skipForwardToDocument}(d)$ 
18:      if  $l_i.\text{getCurrentDocument}() = d$  then
19:         $s_d \leftarrow s_d + g_i(Q)f_i(l_i)$  ▷ Update the document score
20:         $l_i.\text{movePastDocument}( d )$ 
21:      else
22:         $d \leftarrow -1$ 
23:        break
24:      end if
25:    end for
26:    if  $d > -1$  then  $R.\text{add}( s_d, d )$ 
27:    end if
28:  end while
29:  return the top  $k$  results from  $R$ 
30: end procedure

```

Conjunctive Document-at-a-Time


Threshold Methods

- Threshold methods use number of top-ranked documents needed (k) to optimize query processing
 - for most applications, k is small
- For any query, there is a *minimum score* that each document needs to reach before it can be shown to the user
 - score of the k th-highest scoring document
 - gives *threshold* τ
 - optimization methods estimate τ' to ignore documents

Threshold Methods

- For document-at-a-time processing, use score of lowest-ranked document so far for τ'
 - for term-at-a-time, have to use k_{th} -largest score in the accumulator table
- *MaxScore* method compares the maximum score that remaining documents could have to τ'
 - *safe* optimization in that ranking will be the same without optimization

MaxScore Example



- Indexer computes μ_{tree}
 - maximum score for any document containing just “tree”
- Assume $k = 3$, τ' is lowest score after first three docs
- Likely that $\tau' > \mu_{tree}$
 - τ' is the score of a document that contains both query terms
- Can safely skip over all gray postings

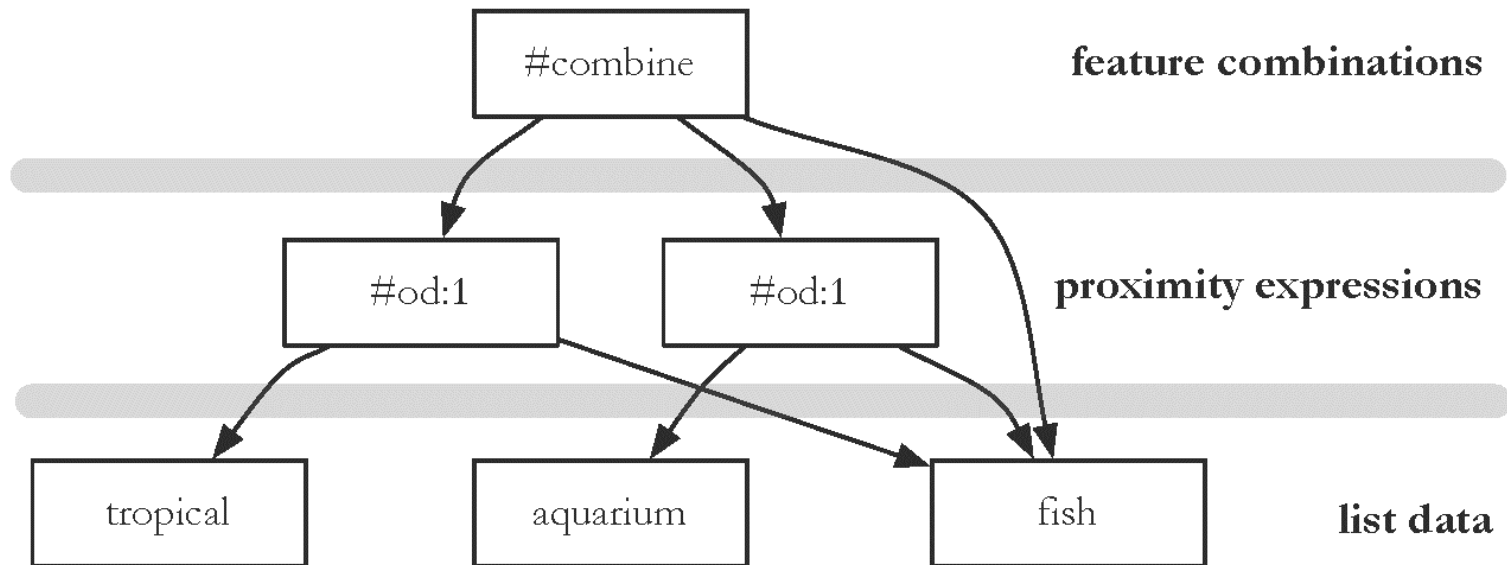
Other Approaches

- Early termination of query processing
 - ignore high-frequency word lists in term-at-a-time
 - ignore documents at end of lists in doc-at-a-time
 - *unsafe* optimization
- List ordering
 - order inverted lists by quality metric (e.g., PageRank) or by partial score
 - makes unsafe (and fast) optimizations more likely to produce good documents

Structured Queries

- *Query language* can support specification of complex features
 - similar to SQL for database systems
 - *query translator* converts the user's input into the structured query representation
 - Galago query language is the example used here
 - e.g., Galago query:
`#combine(#od:1(tropical fish) #od:1(aquarium fish) fish)`

Evaluation Tree for Structured Query



Distributed Evaluation

- Basic process
 - All queries sent to a *director machine*
 - Director then sends messages to many *index servers*
 - Each index server does some portion of the query processing
 - Director organizes the results and returns them to the user
- Two main approaches
 - Document distribution
 - by far the most popular
 - Term distribution

Distributed Evaluation

- Document distribution
 - each index server acts as a search engine for a small fraction of the total collection
 - director sends a copy of the query to each of the index servers, each of which returns the top- k results
 - results are merged into a single ranked list by the director
- Collection statistics should be shared for effective ranking

Distributed Evaluation

- Term distribution
 - Single index is built for the whole cluster of machines
 - Each inverted list in that index is then assigned to one index server
 - in most cases the data to process a query is not stored on a single machine
 - One of the index servers is chosen to process the query
 - usually the one holding the longest inverted list
 - Other index servers send information to that server
 - Final results sent to director

Caching

- Query distributions similar to Zipf
 - About $\frac{1}{2}$ each day are unique, but some are very popular
- Caching can significantly improve effectiveness
 - Cache popular query results
 - Cache common inverted lists
- Inverted list caching can help with unique queries
- Cache must be refreshed to prevent stale data